verMan.io Pitch

What we do

Ingest & Analyze all ML frameworks and AI research

Simulate scenarios against your data & requirements

Output optimization scores against your data

AI industry provides benchmarks against generic data
KitchenSink.ai provides optimizations against your data

Optimization techniques

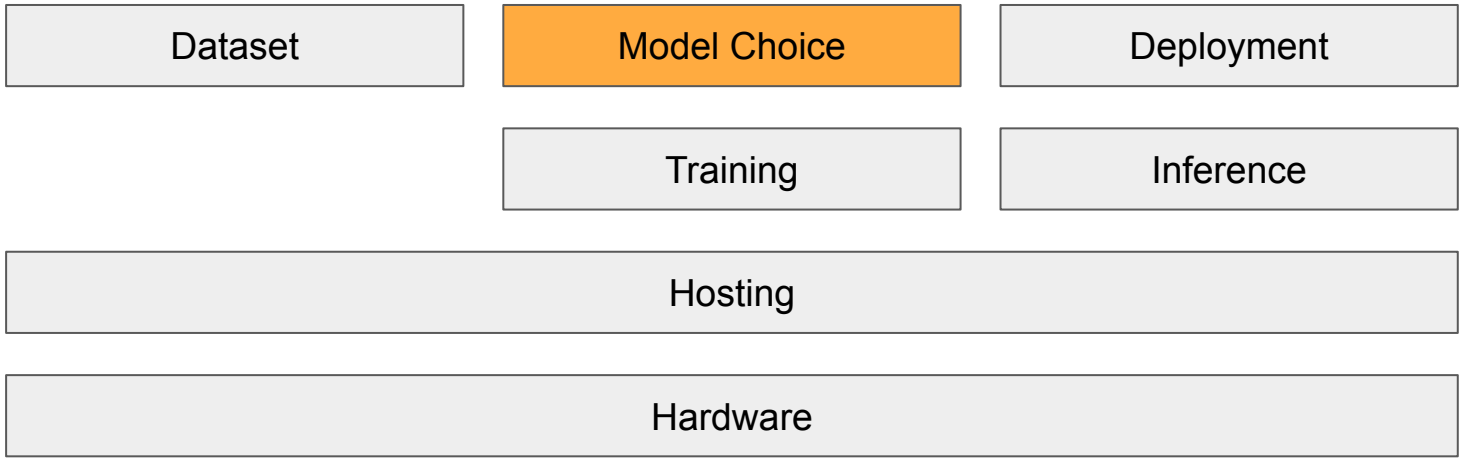

We solve model choice and 3rd-party integrate the rest

Optimization critical to make AI deployment economical

Optimizations can dramatically reduce AI deployment costs

[Need data here]

Endless pipeline of emerging AI research

- [AI state of the art changing rapidly]
- Nearly impossible to analyze and interpret the volume of incoming AI research
- Examples of emerging methodologies []

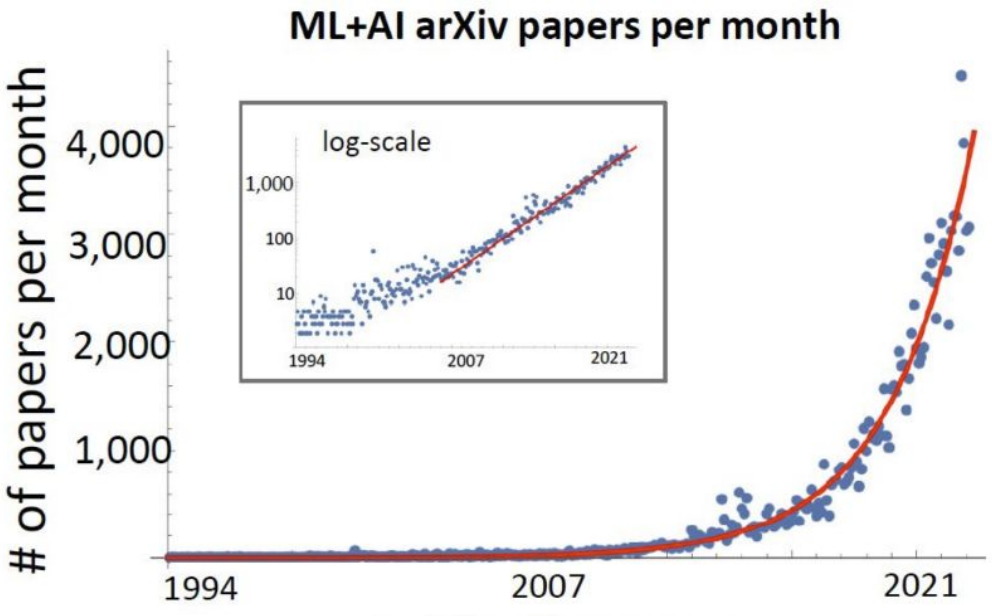

Wide range of optimizations

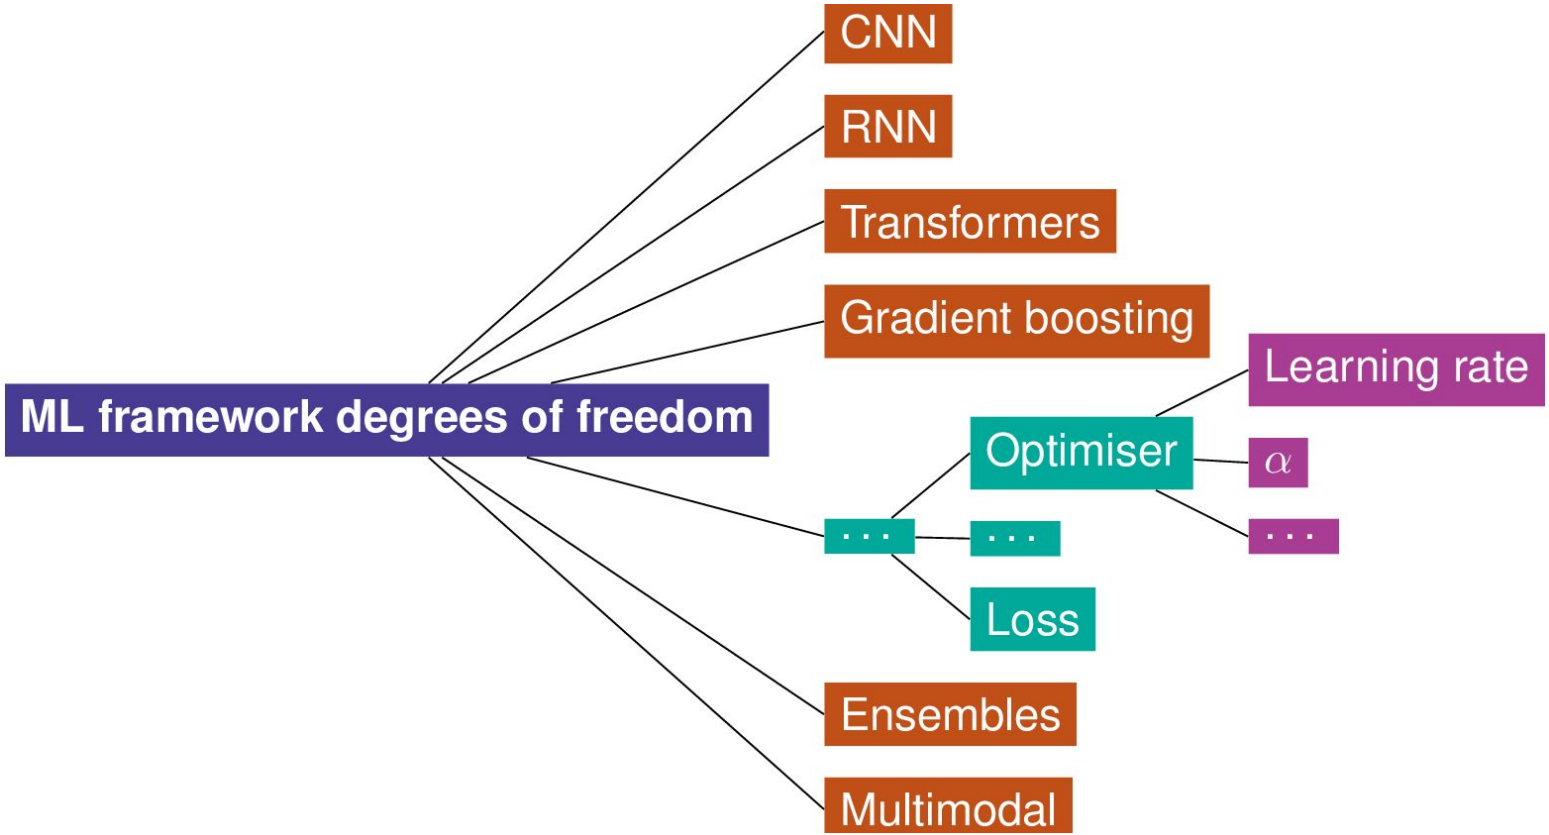

No existing tools allow individualized optimizations at scale

- Economic and performance benefits are unlocked when models optimized against individualized requirements
- Currently, models largely evaluated through leaderboards, which provide a small number of arbitrary benchmarks
- Large organization can run internal bake-offs between models, but this is time consuming

SQL: find best models

```
SELECT model.model_name

FROM Finished_Experiments

WHERE metrics.accuracy > 95.54

  AND dataset == “mnist”

ORDER BY metrics.accuracy;
```

SQL: find cheapest / low-resource models

```
SELECT model.model_name

FROM Finished_Experiments

WHERE metrics.accuracy > 90.54

  AND dataset == “mnist”

  AND model.parameters < 15M

ORDER BY metrics.accuracy;
```

SQL: experiment to find best models

```
SELECT *
FROM Models, Optimizers, Loss
WHERE models IN top_models;
```

^ This output can be fed into a distributed trainer like MLflow and distributed hyperparameter/parameter tuner like Google Vizier, Microsoft NNI, or Ray Tune.

KitchenSink.ai Roadmap

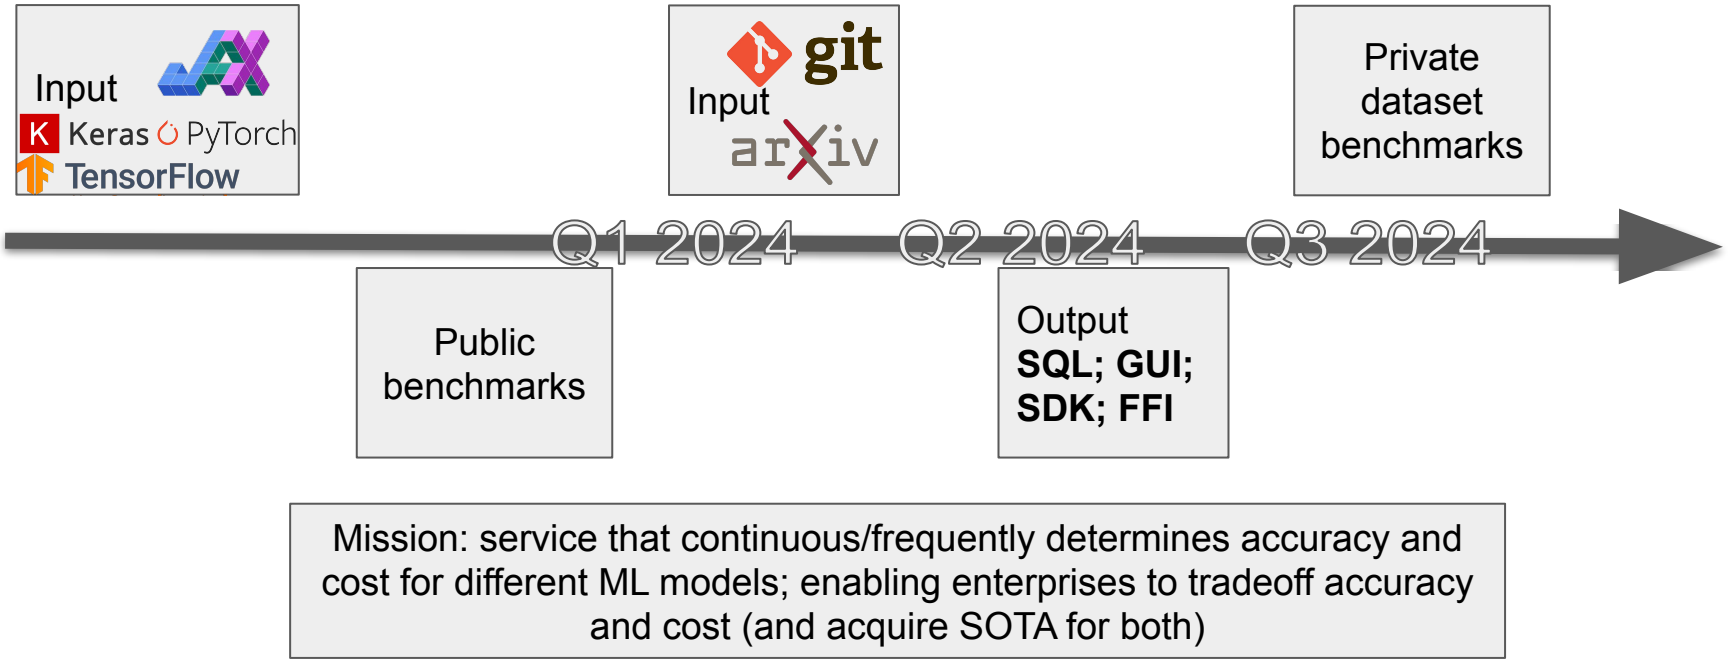

Mission: service that continuous/frequently determines accuracy and cost for different ML models; enabling enterprises to tradeoff accuracy and cost (and acquire SOTA for both)

Team

Samuel Marks, PhD

Finishing off his second PhD in Computer Science (compilers, AI), Samuel was running a software-engineering consultancy for the past decade. He is a top 10 contributor to Keras, and has optimised vector allocations at the C++ level for both TensorFlow & PyTorch. He is also a Google Developer Expert & a Research Fellow at Harvard.

HARVARD MEDICAL SCHOOL

THE UNIVERSITY OF SYDNEY

Sydney and Sydney Eye Hospital

Mass General Brigham
Mass Eye and Ear